

Twitter data serialization methods and benchmarks

Kia Teymourian
Boston University
kiat@bu.edu

Saeed Fathollahzadeh
Free University of
Bozen-Bolzano
s.fathollahzadeh@gmail.com

Chris Jermaine
Rice University
cmj4@rice.edu

ABSTRACT

abstract

PVLDB Reference Format:

Ben Trovato, G. K. M. Tobin, Lars Thørväld, Lawrence P. Leipuner, Sean Fogarty, Charles Palmer, John Smith, Julius P. Kumquat, and Ahmet Sacan. Twitter data serialization methods and benchmarks. *PVLDB*, 12(xxx): xxxx-yyyy, 2019.
DOI: <https://doi.org/10.14778/xxxxxxx.xxxxxxx>

1. INTRODUCTION

Introduction

2. EXPERIMENTAL OVERVIEW

In the next few sections of the paper, we will give detailed explanations of the experimental tasks we consider. As a preview, the tasks we consider are:

1. A set of serialized objects stored externally on an HDD; the task is to read the objects into memory and deserialize them to their in-memory representation.
2. A set of objects are stored in a large file (larger than the available RAM). The task is to perform an external sort of the file in order to perform a duplicate removal.
3. A set of objects are partitioned across a number of machines in a network; the task is to send requests to the machines. Each machine answers the request by serializing the objects, then sending them over the network to the requesting machine.
4. Finally, a set of sparse vectors are stored across various machines on a network. The task is to perform a tree aggregation where the vectors are aggregated over $\log(n)$ hops.

2.1 Twitter Data Set

For the various experiments, we use twitter data sets [1], implemented using each of the ten different physical implementations.

This work is licensed under the Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc-nd/4.0/>. For any use beyond those covered by this license, obtain permission by emailing info@vldb.org. Copyright is held by the owner/author(s). Publication rights licensed to the VLDB Endowment.

Proceedings of the VLDB Endowment, Vol. 12, No. xxx
ISSN 2150-8097.

DOI: <https://doi.org/10.14778/xxxxxxx.xxxxxxx>

Table 1: Frequency of some Tweet Objects (for 1 million tweets)

Object Name	Parent Object	Frequency	AVG Size(byte)
tweet	root object	1,000,000	6599.305602
users	tweet	1,000,000	1453.798722
coordinates	tweet	1586	57.87705
place	tweet	11974	392.8976115
quoted status	tweet	177537	3993.947594
retweeted status	tweet	598517	4775.996797
entities	tweet	1,000,000	336.02128
extended entities	tweet	51479	1133.00505
hashtags	entities	398720	44.0340615
media	entities	51481	762.556516
urls	entities	417176	187.6829563
user mentions	entities	1063216	126.845482
symbols	entities	7793	36.55832
sizes	media	7793	36.55832
media sizes	sizes	51485	194.91743
thumb	media sizes	51485	37.989
large	media sizes	51485	37.524
medium	media sizes	51485	37.412
small	media sizes	51485	36.993

2.2 Encoding sizes

The ten different complex object implementations that we considered have very different encoding densities when the objects are serialized for storage or transmission across the network. The average, per-object sizes are given in Table 1. Figure 1 show histogram of per objects with Standard Deviation(SD).

2.3 Experimental Details

We run our experiments on a virtual server instances which have 8 vCPU cores, 32 GB RAM and one 800 GB hard disks (Network Instance Store) running with Ubuntu 18.04.3 LTS. Before running each experiment task, we “warmed up” the Java Garbage Collector (GC) by creating a large number of objects. We do not include this warm-up-time in our performance time calculations.

We used two Java GC flags `-XX : -UseGCOverheadLimit` and `-XX : +UseConcMarkSweepGC`. The first flag is used to avoid OutOfMemoryError exceptions while using the complete RAM size for data processing and the second flag is for running concurrent garbage collection.

We run all of our experiments 5 times and observed that the results have low variance. In this paper we present the average of those runs. Before running each experiment, we deleted the OS cache using the Linux command: `echo3 >`

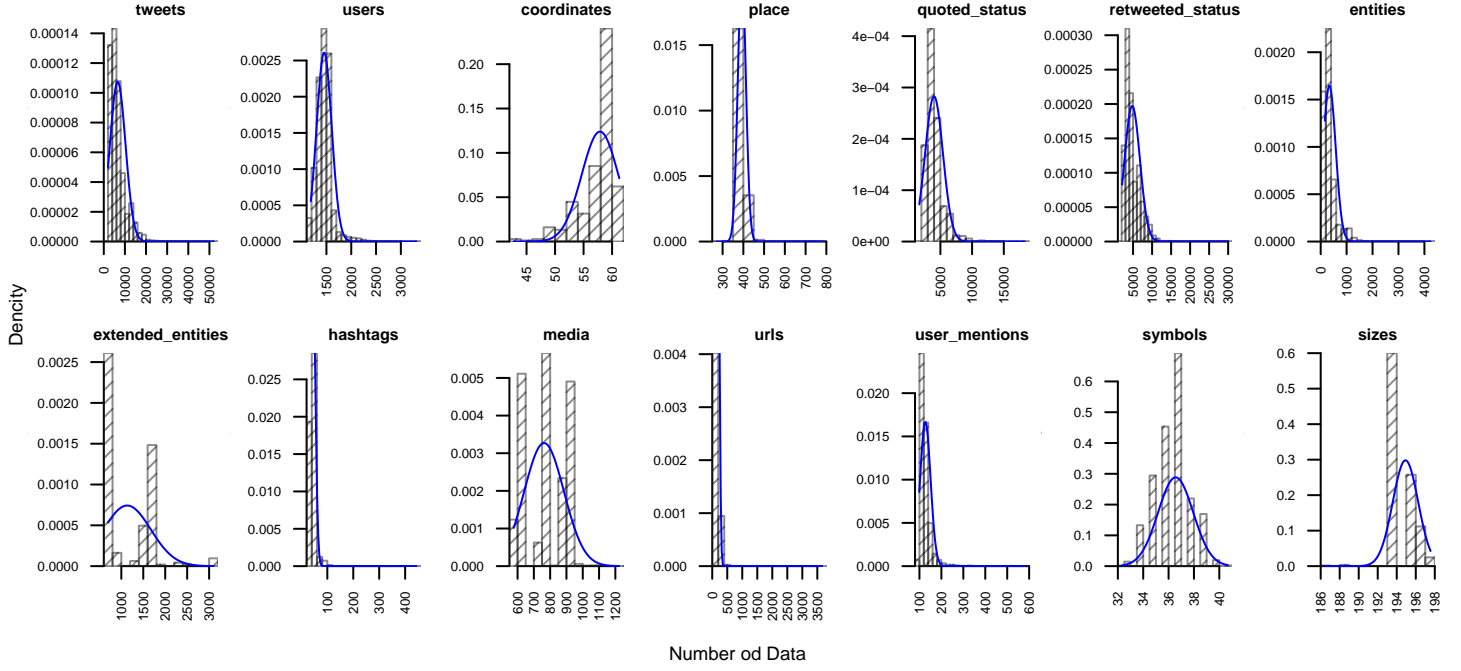


Figure 1: histograms with Standard Deviation

`/proc/sys/vm/drop_caches`.

Our Java implementation is written using Java 8 with the Oracle JDK version "1.8.0.241" and for our C++ implementation we use the C++11, compiled using clang++ (version 6.0.0).

The source codes of our implementation and a brief description of technical details can be found on the Github Repository ¹

3. EXPERIMENTS

¹The source code of our Implementation is available at <https://github.com/fathollahzadeh/serialization>

3.1 I/O FROM LOCAL DISK

The first set of experiments are a set of simple, single-machine experiments. The goal is to examine how the various complex object implementations compare for a simple from-disk retrieval task. In this set of experiments, the part, tweet data set is first loaded onto the HDD drive of a single machine—one version for each of the ten complex object implementations tested—where they are organized into 256KB pages. The objects are then indexed, using a dense index.

Two experiments are run. In the first, a particular object is looked up in the index, and then enough pages are read from disk to access that object, as well as the following $n - 1$ objects. As the pages are loaded into RAM, all n objects are deserialized and made ready for processing. This tests the ability of the object implementation to support fast processing of objects in sequence. We test n in $\{ 1 \times 10^6, 2 \times 10^6, 3 \times 10^6, 4 \times 10^6, 5 \times 10^6 \}$.

In the second experiment, a list of n , randomly-selected objects are accessed, in order. For each object, the location of the object in the database is looked up in the index, and then the corresponding page is loaded into RAM. The desired object is then deserialized from the page. This simulates a scenario where objects are retrieved from secondary storage using a secondary index.

Before the experiment, the operating system buffer cache is emptied. We do not utilize a dedicated buffer cache, but we do allow the operating system to cache disk pages.

One concern is that since there was only one thread active at a given time during each experiment, this might give an unfair advantage to those solutions running in the JVM. One of the classical performance problems observed during garbage collection is long pauses during which worker threads are largely locked out of allocations. When only a single thread is available, serial execution is already forced, and hence the cost of such a pause is minimized. A single-threaded environment might also give an unfair advantage to the non-in-place C++ solutions, as *malloc()* and *free()* may be less of a bottleneck when only a single thread is running. Hence, we also ran a set of sequential read experiments where four threads were allowed to concurrently read and then process different ranges of the required data. This may give a better idea of the performance in a realistic, multi-threaded environment.

3.1.1 Results

3.1.2 Discussion

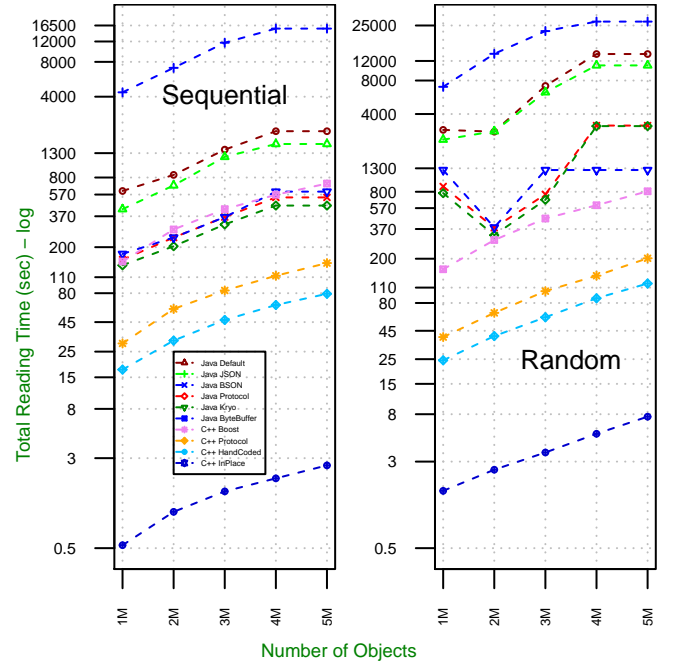


Figure 2: sequential and random read

4. REFERENCES

- [1] Tweet objects.