

Twitter data serialization methods and benchmarks

Saeed Fathollahzadeh
Free University of
Bozen-Bolzano
saeed.fathollahzadeh@unibz.it

Kia Teymourian
Boston University
kiat@bu.edu

Chris Jermaine
Rice University
cmj4@rice.edu

ABSTRACT

abstract

PVLDB Reference Format:

Ben Trovato, G. K. M. Tobin, Lars Thørväld, Lawrence P. Leipuner, Sean Fogarty, Charles Palmer, John Smith, Julius P. Kumquat, and Ahmet Sacan. Twitter data serialization methods and benchmarks. *PVLDB*, 12(xxx): xxxx-yyyy, 2019.
DOI: <https://doi.org/10.14778/xxxxxxx.xxxxxxx>

1. INTRODUCTION

Introduction

The contribution of our work are the following:

1. we implement all serialization methods in C++ and Java programming language in a single thread system. But, we evaluate the methods with *taskset* for restrict run the method on a special core and without *taskset* to allow run method on the arbitrary cores. In the experimental section we mention which methods use thread or which platforms want to improve performance of processing.
2. we implement same method in both C++ and Java programming language and we demonstrated same technique haven't same performance in differ languages(e.g. google protobuf).
3. we compare all methods with a complex data sets. In academic setting over the last decade, there has been significant progress in serialization methods. However, much of this work makes assumptions that are simply unrealistic for deployed industrial applications. In this work, we used twitter dataset. This dataset include more objects type with deep hierarchy. Some methods need to save object meta data in serialization step and will be use it in the de-serialization section.
4. we investigate which methods are easy to used. It means is which methods create transparent view in

This work is licensed under the Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc-nd/4.0/>. For any use beyond those covered by this license, obtain permission by emailing info@vldb.org. Copyright is held by the owner/author(s). Publication rights licensed to the VLDB Endowment.

Proceedings of the VLDB Endowment, Vol. 12, No. xxx
ISSN 2150-8097.

DOI: <https://doi.org/10.14778/xxxxxxx.xxxxxxx>

develop step. For example in C++ *InPlace* we need just one line for de-serialization, But in the serialization step we should spend more times for convert object to the method schema.

5. we evaluate multiple famous serialization methods in big data systems. We focused specially on C++ and Java programming language. So, we deeply compared CPU, Memory and I/O for HDD resources. Our empirical experiments demonstrate best way for choose best method in a big data system.

2. EXPERIMENTAL OVERVIEW

In the next few sections of the paper, we will give detailed explanations of the experimental tasks we consider. As a preview, the tasks we consider are:

1. A set of serialized objects stored externally on an HDD; the task is to read the objects into memory and deserialize them to their in-memory representation.
2. A set of objects are stored in a large file (larger than the available RAM). The task is to perform an external sort of the file in order to perform a duplicate removal.
3. A set of objects are partitioned across a number of machines in a network; the task is to send requests to the machines. Each machine answers the request by serializing the objects, then sending them over the network to the requesting machine.
4. Finally, a set of sparse vectors are stored across various machines on a network. The task is to perform a tree aggregation where the vectors are aggregated over $\log(n)$ hops.

2.1 Twitter Data Set

For the various experiments, we use twitter data sets [1], implemented using each of the ten different physical implementations.

2.2 Encoding sizes

The ten different complex object implementations that we considered have very different encoding densities when the objects are serialized for storage or transmission across the network. The average, per-object sizes are given in Table 3.

Table 1: Frequency of some Tweet Objects (for 1 million tweets)

Object Name	Parent Object	Frequency
tweet	root object	1,000,000
users	tweet	1,000,000
coordinates	tweet	1586
place	tweet	11974
quoted status	tweet	177537
retweeted status	tweet	598517
entities	tweet	1,000,000
extended entities	tweet	51479
hashtags	entities	398720
media	entities	51481
urls	entities	417176
user mentions	entities	1063216
symbols	entities	7793
sizes	media	7793
media sizes	sizes	51485
thumb	media sizes	51485
large	media sizes	51485
medium	media sizes	51485
small	media sizes	51485

Table 2: tweet complexity

Tweet type	Frequency
Simple tweets(retweet & quote are null)	332,901
Retweets	489,562
Quote	68,582
Retweet & Quote	108,955
Total	1,000,000

Table 3: Comparison of object size for 1 million tweet

Serialization Methods	Serialized file size(gigabyte)
Java Default	4.6
Java JSON Gzip	1.4
Java BSON	4.9
Java Protocol Buffer	1.9
Java Kryo	1.9
Java Hand Coded ByteBuffer	2.3
Java FaltBuffers	2.9
C++ Hand Coded	2.1
C++ InPlace	3.2
C++ Boost	2.2
C++ Protocol Buffer	1.9
C++ BSON	4.6
C++ FaltBuffers	2.9

2.3 Experimental Details

We run our experiments on Google Cloud costume instances which have 4 vCPU cores, 32 GB RAM and 3000 GB standard persistent disk(Sustained random IOPS limit: read=2,250 and write=4,500) running with Ubuntu Ubuntu 18.04.4 LTS. Before running each experiment task, we "warmed up" the Java Garbage Collector (GC) by creating a large number of objects. We do not include this warm-up-time in our performance time calculations.

We used two Java GC flags $-XX : -UseGCOverheadLimit$ and $-XX : +UseConcMarkSweepGC$. The first flag is used to avoid OutOfMemoryError exceptions while using the complete RAM size for data processing and the second flag is for running concurrent garbage collection.

We run all of our experiments 3 times and observed that the results have low variance. In this paper we present the average of those runs. Before running each experiment, we deleted th OS cache using the Linux command: `echo3 > /proc/sys/vm/drop_caches`.

Our Java implementation is written using Java 8 with the Oracle JDK version "1.8.0.241" and for our C++ implementation we use the C++11, compiled using clang++ (version 6.0.0).

The source codes of our implementation and a brief description of technical details can be found on the Github Repository ¹.

3. EXPERIMENTS

¹The source code of our Implementation is available at <https://github.com/fathollahzadeh/serialization>




Figure 1: Serialize Objects for 5M Tweets

3.1 Serialize RAW Data to Local Disk

The first step of experiments are serialize various complex objects and the write into a file in disk. In this experiment, the raw tweet data set read line-by-line and convert to a objects. The serialization tasks for each of the thirteen implementation method run. In the serialization process each object serialized or copy the final serialization result into 256KB pages and the objects indexed in separate file.

3.1.1 Results

In Figure 1, we show, for each of the thirteen implementations, for both *taskset* TRUE and FALSE the total running time required as a function of the number of Tweet objects write experiments. In the figure where the performance differences are easier to see; we also breakdown the total time into I/O and CPU.

3.1.2 Discussion

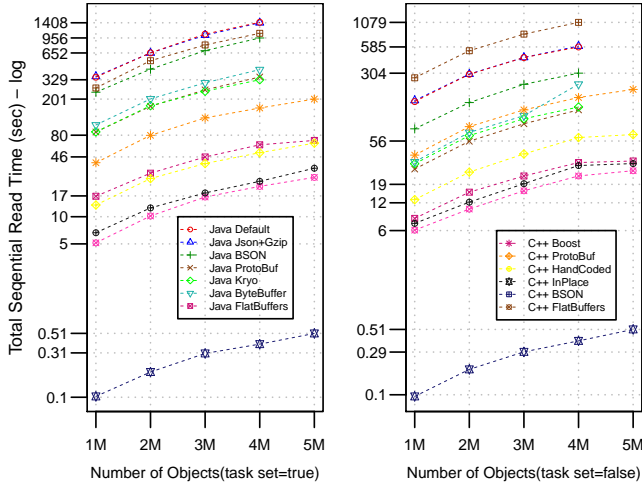


Figure 2: sequential read

3.2 I/O FROM LOCAL DISK

The goal is to examine how the various complex object implementations compare for a simple from-disk retrieval task. In this set of experiments, the tweet data set is first loaded onto the HDD drive of a machine where they are organized into 256KB pages. The objects are then indexed, using a dense index.

Two experiments are run. In the first, a particular object is looked up in the index, and then enough pages are read from disk to access that object, as well as the following $n - 1$ objects. As the pages are loaded into RAM, all n objects are de-serialized and made ready for processing. This tests the ability of the object implementation to support fast processing of objects in sequence. We test n in $\{1 \times 10^6, 2 \times 10^6, 3 \times 10^6, 4 \times 10^6, 5 \times 10^6\}$.

In the second experiment, a list of n , randomly-selected objects are accessed, in order. For each object, the location of the object in the database is looked up in the index, and then the corresponding page is loaded into RAM. The desired object is then de-serialized from the page. This simulates a scenario where objects are retrieved from secondary storage using a secondary index.

Before the experiment, the operating system buffer cache is emptied. We do not utilize a dedicated buffer cache, but we do allow the operating system to cache disk pages.

3.2.1 Results

3.2.2 Discussion




Figure 3: CPU and IO details of sequential read for 4M tweets




Figure 4: random read




Figure 5: CPU and IO details of random read for 4M tweets




Figure 6: memory used in read objects for 4M Tweets

3.3 Exp. Memory usage

4. REFERENCES

- [1] Tweet objects.